

```
!pip install transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.11/dist-packages (4.51.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from transformers) (3.18.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.30.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.30.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (24.1)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (6.0.2)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2024.11.23)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from transformers) (2.32.3)
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.21.0)
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.5.2)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.11/dist-packages (from transformers) (4.67.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0,>=0.30.0) (2025.5.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0,>=0.30.0) (4.12.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests) (3.10.1)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests) (2025.11.11)
```

Double-click (or enter) to edit

```
from transformers import pipeline
```

```
a = pipeline("sentiment-analysis")
```

```
Based on the provided information, the model used is distilbert-base-uncased-finetuned-sst-2-english and revision 714eb0f (https://huggingface.co/distilbert/distilbert-base-uncased-finetuned-sst-2-english).
Production is not recommended.
```

```
text = input("Enter the text: ")
print(a(text))
```

```
Enter the text: i love this product. its amazing and works perfectly
[{'label': 'POSITIVE', 'score': 0.9998792409896851}]
```

```
gen = pipeline('text-generation')
```

```
No model was supplied, defaulted to openai-community/gpt2 and revision 607a30d (https://huggingface.co/openai-community/gpt2).
Using a pipeline without specifying a model name and revision in production is not recommended.
```

```
config.json: 100% 665/665 [00:00<00:00, 9.92kB/s]
```

```
Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed.
```

```
model.safetensors: 100% 548M/548M [00:07<00:00, 25.2MB/s]
```

```
generation_config.json: 100% 124/124 [00:00<00:00, 10.5kB/s]
```

```
tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 2.19kB/s]
```

```
vocab.json: 100% 1.04M/1.04M [00:00<00:00, 5.51MB/s]
```

```
merges.txt: 100% 456k/456k [00:00<00:00, 19.9MB/s]
```

```
tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 17.1MB/s]
```

```
Device set to use cpu
```

```
text = input("Enter the text: ")
t = gen(text, max_length = 200)
```

Enter the text: once upon a time there lived a king mufasa
 Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.

```
print(t[0].get('generated_text'))
```

i who sent his son to marry the daughter of an abbe. She was not brought to see her father, and so she was unable to

```
n = pipeline('ner')
```

No model was supplied, defaulted to dbmdz/bert-large-cased-finetuned-conll03-english and revision 4c53496 (<https://huggingface.co/dbmdz/bert-large-cased-finetuned-conll03-english>)
 Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100%

998/998 [00:00<00:00, 15.9kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download
 WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed

model.safetensors: 100%

1.33G/1.33G [00:14<00:00, 224MB/s]

Some weights of the model checkpoint at dbmdz/bert-large-cased-finetuned-conll03-english were not used when initializing the model
 - This is expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task (e.g. sequence classification) as they have different vocabularies
 - This is NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to be initialized with those weights

tokenizer_config.json: 100%

60.0/60.0 [00:00<00:00, 4.76kB/s]

vocab.txt: 100%

213k/213k [00:00<00:00, 1.69MB/s]

Device set to use cpu


```
text = input("Enter a text: ")
result = n(text)
result
```

Enter a text: i am dhilleswari from srikakulam

```
[{'entity': 'I-PER',
  'score': np.float32(0.81985587),
  'index': 3,
  'word': 'd',
  'start': 5,
  'end': 6},
 {'entity': 'I-PER',
  'score': np.float32(0.67239696),
  'index': 4,
  'word': '##hill',
  'start': 6,
  'end': 10},
 {'entity': 'I-LOC',
  'score': np.float32(0.68562746),
  'index': 8,
  'word': 's',
  'start': 22,
  'end': 23},
 {'entity': 'I-LOC',
  'score': np.float32(0.45079637),
  'index': 9,
  'word': '##rika',
  'start': 23,
  'end': 27},
 {'entity': 'I-LOC',
  'score': np.float32(0.50737554),
  'index': 10,
  'word': '##ku',
  'start': 27,
  'end': 29},
 {'entity': 'I-LOC',
  'score': np.float32(0.45141292),
  'index': 11,
  'word': '##lam',
  'start': 29,
  'end': 33}]
```

```
'start': 29,
'end': 32}]
```

```
qa = pipeline('question-answering')
```

 No model was supplied, defaulted to distilbert/distilbert-base-cased-distilled-squad and revision 564e9b5 (<https://huggingface.co/distilbert/distilbert-base-cased-distilled-squad>)
Using a pipeline without specifying a model name and revision in production is not recommended.

```
config.json: 100% 473/473 [00:00<00:00, 7.50kB/s]
```

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed

```
model.safetensors: 100% 261M/261M [00:02<00:00, 121MB/s]
```

```
tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 4.43kB/s]
```

```
vocab.txt: 100% 213k/213k [00:00<00:00, 7.16MB/s]
```


```
tokenizer.json: 100% 436k/436k [00:00<00:00, 3.28MB/s]
```

```
Device set to use cpu
```

```
c = '''bity bought some buter, but the buter was bitter butter.So, Bittey bought some better butter,to make bitter but
```

```
q = "what does Bittey bought?"
```


```
result = qa({'question':q, 'context': c})
print(result)
```

 {'score': 0.33088603615760803, 'start': 12, 'end': 22, 'answer': 'some buter'}
/usr/local/lib/python3.11/dist-packages/transformers/pipelines/question_answering.py:390: FutureWarning: Passing a list to warnings.warn() is more deprecation than the warning message
warnings.warn(

```
c = '''The sun rises in the east and sets in the west. This daily cycle is a fundamental aspect of life on Earth, infl
```

```
q = "What role does sunlight play in the process of photosynthesis?"
```

```
result = qa({'question':q, 'context': c})
print(result)
```

 /usr/local/lib/python3.11/dist-packages/transformers/pipelines/question_answering.py:390: FutureWarning: Passing a list to warnings.warn() is more deprecation than the warning message
warnings.warn(
{'score': 0.6297405362129211, 'start': 218, 'end': 224, 'answer': 'energy'}

```
s = pipeline('summarization')
```

➡ No model was supplied, defaulted to sshleifer/distilbart-cnn-12-6 and revision a4f8f3e (<https://huggingface.co/sshleifer/distilbart-cnn-12-6>)
Using a pipeline without specifying a model name and revision in production is not recommended.

```
config.json: 100% 1.80k/1.80k [00:00<00:00, 60.6kB/s]
pytorch_model.bin: 100% 1.22G/1.22G [00:16<00:00, 50.7MB/s]
model.safetensors: 21% 262M/1.22G [00:11<00:18, 52.5MB/s]
tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 471B/s]
vocab.json: 100% 899k/899k [00:00<00:00, 7.24MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 3.46MB/s]
Device set to use cpu
```

```
t = '''Computer vision is a field of artificial intelligence (AI)
that focuses on enabling computers to "see" and interpret digital images and videos in a way that humans do.
It involves developing algorithms and techniques that allow machines to process visual data, extract meaningful information
and perform tasks like object recognition, image classification, and scene understanding.'''
```

```
result = s(t)
result[0].get('summary_text')
```

➡ Your max_length is set to 142, but your input_length is only 72. Since this is a summarization task, where outputs are longer than inputs, this is not recommended. It involves developing algorithms and techniques that allow machines to process visual data, extract meaningful information, and perform tasks like object recognition, image classification, and scene understanding.

```
u = pipeline('fill-mask')
```

➡ No model was supplied, defaulted to distilbert/distilroberta-base and revision fb53ab8 (<https://huggingface.co/distilbert/distilroberta-base>)
Using a pipeline without specifying a model name and revision in production is not recommended.

```
config.json: 100% 480/480 [00:00<00:00, 8.28kB/s]
Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed
model.safetensors: 100% 331M/331M [00:07<00:00, 13.4MB/s]
Some weights of the model checkpoint at distilbert/distilroberta-base were not used when initializing RobertaForMaskedLM
- This is expected if you are initializing RobertaForMaskedLM from the checkpoint of a model trained on another task
- This IS NOT expected if you are initializing RobertaForMaskedLM from the checkpoint of a model that you expect to be trained on this task
tokenizer_config.json: 100% 25.0/25.0 [00:00<00:00, 1.85kB/s]
vocab.json: 100% 899k/899k [00:00<00:00, 4.60MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 3.46MB/s]
tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 18.4MB/s]
Device set to use cpu
```

```
result = u("I'm teaching <mask> <mask> board")
result
```

➡

```
[{'score': 0.02794056199491024,
 'token': 127,
 'token_str': ' my',
 'sequence': '<s>I'm teaching my<mask> board</s>'},
 {'score': 0.020591234788298607,
 'token': 2185,
 'token_str': ' myself',
 'sequence': '<s>I'm teaching myself<mask> board</s>'},
 {'score': 0.017040545120835304,
```

```

'token': 5,
'token_str': ' the',
'sequence': "<s>I'm teaching the<mask> board</s>",
{'score': 0.015600067563354969,
'token': 1159,
'token_str': ' kids',
'sequence': "<s>I'm teaching kids<mask> board</s>"},
{'score': 0.015396728180348873,
'token': 10,
'token_str': ' a',
'sequence': "<s>I'm teaching a<mask> board</s>"}],
[{'score': 0.21203042566776276,
'token': 31960,
'token_str': ' bulletin',
'sequence': "<s>I'm teaching<mask> bulletin board</s>"},
{'score': 0.14683669805526733,
'token': 12733,
'token_str': ' ija',
'sequence': "<s>I'm teaching<mask> ija board</s>"},
{'score': 0.08063647150993347,
'token': 15,
'token_str': ' on',
'sequence': "<s>I'm teaching<mask> on board</s>"},
{'score': 0.07296892255544662,
'token': 5523,
'token_str': ' drawing',
'sequence': "<s>I'm teaching<mask> drawing board</s>"},
{'score': 0.03655911982059479,
'token': 9326,
'token_str': ' circuit',
'sequence': "<s>I'm teaching<mask> circuit board</s>"}]]

```

Start coding or [generate](#) with AI.